

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
```

```
In [2]: data = pd.read_csv('hou_all.csv')
```

```
In [3]: data.head()
```

```
Out[3]:
```

	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	black	lstat	medv	Ur
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	15.3	396.90	4.98	24.0	1
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	396.90	9.14	21.6	1
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242	17.8	392.83	4.03	34.7	1
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7	394.63	2.94	33.4	1
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7	396.90	5.33	36.2	1

```
In [4]: data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 15 columns):
crim      506 non-null float64
zn        506 non-null float64
indus     506 non-null float64
chas      506 non-null int64
nox       506 non-null float64
rm        506 non-null float64
age       506 non-null float64
dis       506 non-null float64
rad       506 non-null int64
tax       506 non-null int64
ptratio   506 non-null float64
black     506 non-null float64
lstat     506 non-null float64
medv     506 non-null float64
Unnamed: 14  506 non-null int64
dtypes: float64(11), int64(4)
memory usage: 59.4 KB
```

```
In [6]: data.isna().sum()
```

```
Out[6]: crim      0
zn      0
indus    0
chas     0
nox      0
rm       0
age      0
dis      0
rad      0
tax      0
ptratio  0
black    0
lstat    0
medv     0
Unnamed: 14  0
dtype: int64
```

Last Column Unnamed: 14 doesn't have and significance so remove it

```
In [8]: data = data.drop('Unnamed: 14',axis=1)
```

```
In [9]: data.head()
```

Out[9]:

	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	black	lstat	medv
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	15.3	396.90	4.98	24.0
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	396.90	9.14	21.6
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242	17.8	392.83	4.03	34.7
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7	394.63	2.94	33.4
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7	396.90	5.33	36.2

Generally, NaN or missing values can be in any form like 0, ? or may be written as "missing" and in our case, we can see that there are a lot of 0's, so we can replace them with NaN to calculate how much data we are missing.

```
In [10]: data.zn.replace(0,np.nan,inplace = True)
data.chas.replace(0,np.nan,inplace=True)
```

```
In [11]: data.head()
```

Out[11]:

	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	black	lstat	medv
0	0.00632	18.0	2.31	NaN	0.538	6.575	65.2	4.0900	1	296	15.3	396.90	4.98	24.0
1	0.02731	NaN	7.07	NaN	0.469	6.421	78.9	4.9671	2	242	17.8	396.90	9.14	21.6
2	0.02729	NaN	7.07	NaN	0.469	7.185	61.1	4.9671	2	242	17.8	392.83	4.03	34.7
3	0.03237	NaN	2.18	NaN	0.458	6.998	45.8	6.0622	3	222	18.7	394.63	2.94	33.4
4	0.06905	NaN	2.18	NaN	0.458	7.147	54.2	6.0622	3	222	18.7	396.90	5.33	36.2

```
In [12]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 14 columns):
crim      506 non-null float64
zn        134 non-null float64
indus     506 non-null float64
chas      35 non-null float64
nox       506 non-null float64
rm        506 non-null float64
age       506 non-null float64
dis       506 non-null float64
rad       506 non-null int64
tax       506 non-null int64
ptratio   506 non-null float64
black     506 non-null float64
lstat     506 non-null float64
medv      506 non-null float64
dtypes: float64(12), int64(2)
memory usage: 55.4 KB
```

```
In [13]: #Percent of data which is not available
data.isnull().sum()/len(data) * 100
```

```
Out[13]: crim      0.000000
zn          73.517787
indus      0.000000
chas      93.083004
nox        0.000000
rm         0.000000
age        0.000000
dis        0.000000
rad        0.000000
tax        0.000000
ptratio    0.000000
black      0.000000
lstat      0.000000
medv      0.000000
dtype: float64
```

both "ZN" and "CHAS" are missing more than 70% data, so will remove both these features

```
In [14]: data = data.drop(['zn','chas'],axis=1)
```

```
In [16]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 12 columns):
crim      506 non-null float64
indus     506 non-null float64
nox       506 non-null float64
rm        506 non-null float64
age       506 non-null float64
dis       506 non-null float64
rad       506 non-null int64
tax       506 non-null int64
ptratio   506 non-null float64
black     506 non-null float64
lstat     506 non-null float64
medv     506 non-null float64
dtypes: float64(10), int64(2)
memory usage: 47.5 KB
```

```
In [17]: data.describe()
```

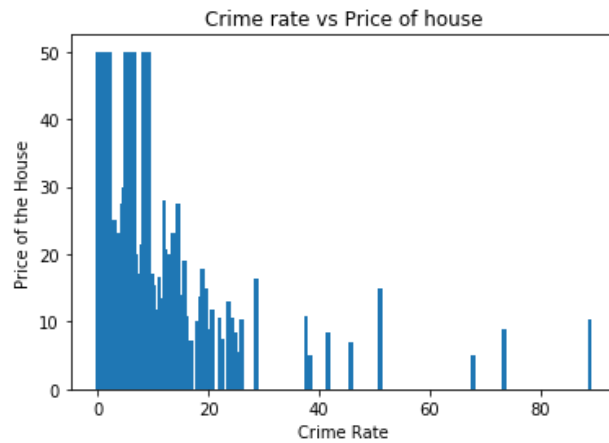
```
Out[17]:
```

	crim	indus	nox	rm	age	dis	rad	
<b>count</b>	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	50
<b>mean</b>	3.613524	11.136779	0.554695	6.284634	68.574901	3.795043	9.549407	40
<b>std</b>	8.601545	6.860353	0.115878	0.702617	28.148861	2.105710	8.707259	16
<b>min</b>	0.006320	0.460000	0.385000	3.561000	2.900000	1.129600	1.000000	18
<b>25%</b>	0.082045	5.190000	0.449000	5.885500	45.025000	2.100175	4.000000	27
<b>50%</b>	0.256510	9.690000	0.538000	6.208500	77.500000	3.207450	5.000000	33
<b>75%</b>	3.677082	18.100000	0.624000	6.623500	94.075000	5.188425	24.000000	66
<b>max</b>	88.976200	27.740000	0.871000	8.780000	100.000000	12.126500	24.000000	71

### Separating the dependent and independent variables

```
In [21]: plt.bar(data.crim,data.medv)
plt.xlabel('Crime Rate')
plt.ylabel('Price of the House')
plt.title('Crime rate vs Price of house')
plt.plot()
```

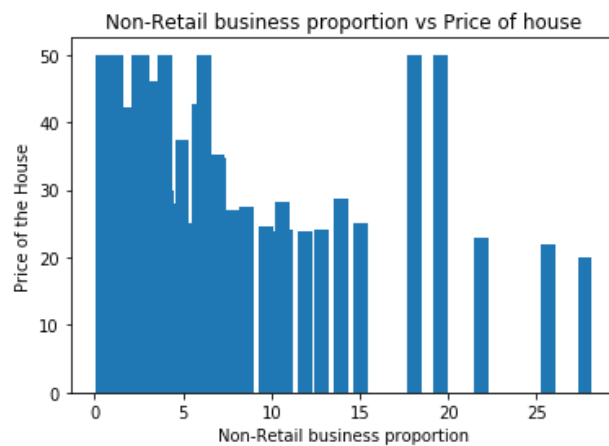
Out[21]: []



### As Crime rate increases the rate of House decreases


```
In [22]: plt.bar(data.indus,data.medv)
plt.xlabel('Non-Retail business proportion')
plt.ylabel('Price of the House')
plt.title('Non-Retail business proportion vs Price of house')
plt.plot()
```

Out[22]: []



```
In [23]: plt.bar(data.nox,data.medv)
plt.xlabel('Nitric Oxide concentration')
plt.ylabel('Price of the House')
plt.title('Nitric Oxide concentration vs Price of house')
plt.plot()
```

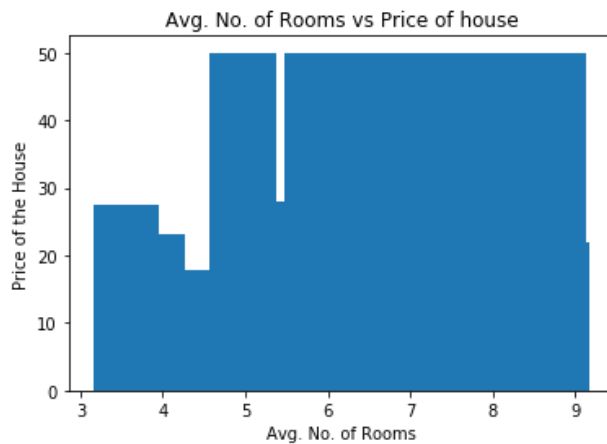
Out[23]: []



**As Nitric Oxide concentration increases the rate of House decreases**

```
In [24]: plt.bar(data.rm,data.medv)
plt.xlabel('Avg. No. of Rooms')
plt.ylabel('Price of the House')
plt.title('Avg. No. of Rooms vs Price of house')
plt.plot()
```

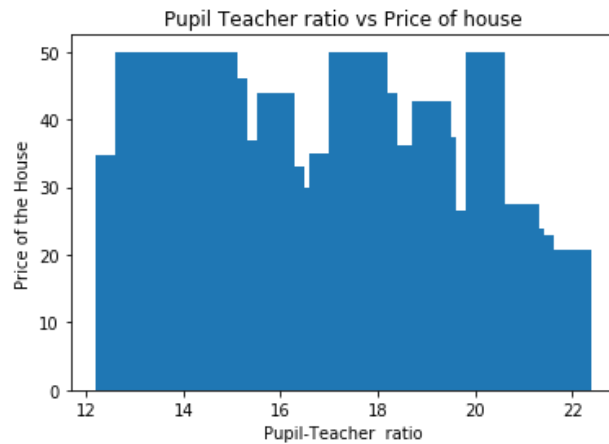
Out[24]: []



**As no. of rooms increases price of house increases**

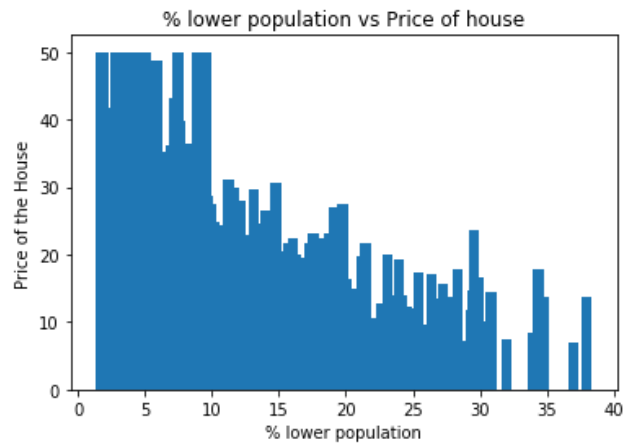
```
In [25]: plt.bar(data.ptratio,data.medv)
plt.xlabel('Pupil-Teacher ratio')
plt.ylabel('Price of the House')
plt.title('Pupil Teacher ratio vs Price of house')
plt.plot()
```

Out[25]: []



```
In [36]: plt.bar(data.lstat,data.medv)
plt.xlabel('% lower population')
plt.ylabel('Price of the House')
plt.title('% lower population vs Price of house')
plt.plot()
```

Out[36]: []



Where the lower status population is low price of houses are high

```
In [27]: X = data.iloc[:, :-1]
y = data.iloc[:, -1]
```

```
In [30]: from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.30,random_state=0)
```

```
In [33]: from sklearn.linear_model import LinearRegression
model = LinearRegression()
model.fit(X_train, y_train) # Fitting our model to the training set

Out[33]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1, normalize=False)

In [34]: y_pred = model.predict(X_test)
```


#### Print the coefficient of the Model

```
In [37]: # The coefficients
print('Coefficients: \n', model.coef_)

Coefficients:
[-1.19332593e-01  1.20970593e-02 -1.73005334e+01  4.12170424e+00
 -1.37417904e-02 -1.25055232e+00  2.44921612e-01 -9.67542136e-03
 -1.18932063e+00  7.56361429e-03 -4.79845036e-01]
```

```
In [38]: plt.scatter(y_test, y_pred)
plt.xlabel('Y Test')
plt.ylabel('Predicted Y')
```

```
Out[38]: Text(0,0.5,'Predicted Y')
```



```
In [39]: from sklearn import metrics
print("MAE", metrics.mean_absolute_error(y_test, y_pred))
print("MSE", metrics.mean_squared_error(y_test, y_pred))
print("RMSE", np.sqrt(metrics.mean_squared_error(y_test, y_pred)))
print("Score:", model.score(X_test, y_test))

MAE 3.7125347264819153
MSE 28.477481110580698
RMSE 5.336429622002027
Score: 0.6579918387799596
```